

HEART: HiErarchical Abstraction for Real-Time Partial Order Planning

Paper #1788

Abstract

When asked about the biggest issues with automated planning, the answer often consists of speed, quality, and expressivity. However, certain applications need a mix of these criteria. Most existing solutions focus on one area while not addressing the others. We aim to present a new method to compromise on these aspects in respect to demanding domains like assistive robotics, intent recognition, and real-time computations. The HEART planner is an anytime planner based on a hierarchical version of Partial Order Planning (POP). Its principle is to retain a partial plan at each abstraction level in the hierarchy of actions. While the intermediary plans are not solutions, they meet criteria for usages ranging from explanation to goal inference. This paper also evaluates the variations of speed/quality combinations of the abstract plans in relation to the complete planning process.

Introduction

One of the most discussed issues in the domain of assistive robotics is intent recognition. Recently some new methods emerged to tackle this problem (Talamadupula *et al.* 2014). However, the domain requires being able to solve expensive human-related problems on robots with limited processing power and to provide real-time answers, often several times per second.

Planners are not meant to solve intent recognition problems. However, several works extended what is called in psychology the *theory of mind*. That theory is the equivalent of asking “*what would I do if I was them ?*” when observing the behavior of other agents. This leads to new ways to use *inverted planning* as an inference tool. One of the first to propose that idea was Baker *et al.* (2007) that use Bayesian planning to infer intentions. Ramirez and Geffner (2009) found an elegant way to transform a plan recognition problem into classical planning. This is done simply by encoding observation constraints in the planning domain (Baiocchi *et al.* 1998) to ensure the selection of actions in the order they were observed. A cost comparison will then give a probability of the goal to be pursued given the observations. Chen *et al.* (2013) extended this with multi-goal recognition. A new method, proposed by Sohrabi *et al.* (2016), makes the recognition fluent centric. It assigns costs to missing

or noisy observed fluents, which allows finer details and less preprocessing work than action based recognition. This method also uses a meta-goal that combines each possible goal and is realized when at least one of these goals is satisfied. Sohrabi *et al.* state that the quality of the recognition is directly linked to the properties of the generated plans. Thus guided diverse planning¹ was preferred along with the ability to infer several probable goals at once.

The impact of problems size on performance is one of the main issues of automated planning as it has been proven to be a *P-SPACE* problem if not harder (Weld 1994). Sizable problems are usually intractable within the tight constraints of real-time robotics. The problem of dealing with unsolvability has already been addressed by Göbelbecker *et al.* (2010) where “excuses” are being investigated as potential explanations for when a problem has no solutions. The closest way to address time unsolvability is by using explainability (Fox *et al.* 2017). In our context, we do not need the complete solution as much as an explanation of the plan that can be exploited to derive probabilities.

For intent recognition, an important metric is the number of correct fluents. So if finding a complete solution is impossible, a partial solution can meet enough criteria to give a good approximation of the goal probability. One of the main approaches to automated planning is called **Partial Order Planning (POP)**. It works by refining partial plans into a solution. The process and resulting plans are more flexible than classical linear ones. Another approach is **Hierarchical Task Networks (HTN)** that are meant to tackle the problem using composite actions in order to define hierarchical tasks within the plan. These two approaches are not exclusive and have been combined in previous works (Gerevini *et al.* 2008; Pellier and Fiorino 2007). Our work is based on **Hierarchical POP (HiPOP)** by Bechon *et al.* (2014). The idea is to expand the classical POP algorithm with new flaws in order to make it compatible with HTN problems. Planning with HTN results in the possibility of extracting high-level plans, useful when applied to intent recognition (Höller *et al.* 2018).

In the rest of the paper, we detail how HiErarchical Abstraction for Real-Time partial order planning (HEART) creates abstract intermediary plans that can be used for intent recognition within

¹Diverse planning is set to find a set of m plans that are distant of d from one another.

Table 1: Our notations are adapted from (Ghallab *et al.* 2004). The symbol \pm demonstrates when the notation has signed variants.

Symbol	Description
\mathcal{D}, \mathcal{P}	Planning domain and problem.
$pre(a), eff(a)$	Preconditions and effects of the action a .
$methods(a)$	Methods of the action a .
$\phi^\pm(l)$	Signed incidence function for partial order plans. ϕ^- gives the source and ϕ^+ the target step of l . No sign gives a pair corresponding to link l .
$L^\pm(a)$	Set of incoming (L^-) and outgoing (L^+) links of step a . No sign gives all adjacent links.
$a_s \xrightarrow{c} a_t$	Link with source a_s , target a_t and cause c .
$causes(l)$	Gives the causes of a causal link l .
$a_a > a_s$	A step a_a is anterior to the step a_s .
A_x^n	Proper actions set of x down n levels. A_x for $n = 1$ and A_x^* for $n = lv(x)$.
$lv(x)$	Abstraction level of the entity x .
$a \triangleright^\pm a'$	Transpose the links of action a onto a' .
$l \downarrow a$	Link l participates in the partial support of step a .
$\pi \downarrow a$	Plan π fully supports a .
$\dagger_f a$	Subgoal : Fluent f isn't supported in step a .
$a_b \boxtimes l$	Threat : Breaker action a_b threatens causal link l .
$a \oplus_m$	Expansion of composite action a using method m .
$var : exp$	The colon is a separator to be read as "such that".
$[exp]$	Iverson's brackets: 0 if $exp = false$, 1 otherwise.

our time constraints. First, we present some definitions for POP, then we explain our method and prove its properties to finally discuss the experimental results.

Definitions

Planners often work in two phases: first we compile the planning domain then we give the planner an instance of a corresponding planning problem to solve.

Domain

The domain specifies the allowed operators that can be used to plan and all the fluents they use as preconditions and effects.

Definition 1 (Domain). A domain is a tuple $\mathcal{D} = \langle C_{\mathcal{D}}, R, F, O \rangle$ where:

- $C_{\mathcal{D}}$ is the set of **domain constants**.
- R is the set of **relations** (also called *properties*) of the domain. These relations are similar to quantified predicates in first order logic.
- F is the set of **fluents** used in the domain to describe operators.
- O is the set of **operators** which are fully lifted *actions*.

Example: The example domain in listing 1 is inspired from the kitchen domain of Ramirez and Geffner (2010).

```
1 take(item) pre (taken(~), ?(item));
  //?(item) is used to make item into
  a variable.
```

```
2 take(item) eff (taken(item));
3 heat(thing) pre (~(hot(thing)),
  taken(thing));
4 heat(thing) eff (hot(thing));
5 pour(thing, into) pre (thing ~(in) into,
  taken(thing));
6 pour(thing, into) eff (thing in into);
7 put(utensil) pre (~(placed(utensil)),
  taken(utensil));
8 put(utensil) eff (placed(utensil),
  ~(taken(utensil)));
9 infuse(extract, liquid, container) ::
  Action; //Composite action of level 1
10 infuse(extract, liquid, container) method (
11
12   init(infuse(extract, liquid, container))
13   -> take(extract),
14
15   init(infuse(extract, liquid, container))
16   -> take(liquid),
17   take(liquid) -> heat(liquid),
18   heat(liquid) -> pour(liquid,
19   container),
20   take(extract) -> pour(extract,
21   container),
22   pour(liquid, container) ->
23   goal(infuse(extract, liquid, container)),
24   pour(extract, container) ->
25   goal(infuse(extract, liquid, container)),
26   heat(liquid) ->
27   goal(infuse(extract, liquid, container))
28 );
29 water :: Liquid; // To prevent these
  entities to become variables
30 cup :: Container;
31 spoon :: Ustensil;
32 make(drink) :: Action; // Level 2
33 make(drink) method (
34   init(make(drink)) -> take(spoon),
35   take(spoon) -> put(spoon),
36   init(make(drink)) ->
37   infuse(drink, water, cup),
38   infuse(drink, water, cup) -> take(cup),
39   take(cup) -> put(cup),
40   put(spoon) -> goal(make(drink)),
41   infuse(drink, water, cup) ->
42   goal(make(drink)),
43   put(cup) -> goal(make(drink))
44 );
```

Listing 1: Domain file used in our planner.

Definition 2 (Fluent). A fluent f is a parameterized boolean relation $r(arg_1, arg_2, \dots, arg_n)$ where:

- $arg_{i \in [1, n]}$ are the arguments (possibly quantified).
- $n = |r|$ is the arity of r .

Fluents are signed. Negative fluents are noted $\neg f$ and behave as a logical complement. The quantifiers are affected by the sign of the fluents. We do not use the closed world hypothesis: fluents are not satisfied until provided whatever their sign. Sets of fluents have a boolean value that equals to the conjunction of all its fluents.

Example: To describe an item not being held we use the fluent $\neg taken(item)$. If the cup contains water $in(water, cup)$ is true.

The central notion of planning is operators. Instantiated operators are usually called *actions*. In our framework, actions can be partially instantiated. We use the term action for both lifted and grounded operators.

Definition 3 (Action). An action is a parametrized tuple $a(args) = \langle name, pre, eff, methods \rangle$ where:

- *name* is the **name** of the action.
- *pre* and *eff* are sets of fluents that are respectively the **pre-conditions** and the **effects** of the action.
- *methods* is a set of **methods** (partial order plans) that can realize the action. Methods, and the methods of their enclosed actions, cannot contain the parent action.

Example: The precondition of the operator $take(item)$ is simply a single negative fluent noted $\neg taken(item)$ ensuring the variable *item* isn't already taken.

Composite actions are represented using methods. An action without methods is called *atomic*. Methods are plans used only in HTN variants of POP.

Definition 4 (Plan). A partially ordered plan is an *acyclic* directed graph $\pi = (S, L)$, with:

- *S* the set of **steps** of the plan as vertices. A step is an action belonging in the plan. *S* must contain an initial step I_π and goal step G_π .
- *L* the set of **causal links** of the plan as edges. We note $l = a_s \xrightarrow{c} a_t$ the link between its source a_s and its target a_t caused by the set of fluent *c*.

In HEART, *ordering constraints* are defined as the transitive cover of causal links over the set of steps. We note ordering constraints: $a_a > a_s$, with a_a being *anterior* to its *successor* a_s . Ordering constraints can't form cycles, meaning that the steps must be different and that the successor can't also be anterior to its anterior steps: $a_a \neq a_s \wedge a_s \not> a_a$. In all plans, the initial and goal steps have their order guaranteed: $I_\pi > G_\pi \wedge \nexists a_x \in S_\pi : a_x > I_\pi \vee G_\pi > a_x$. If we need to enforce order, we simply add a causal link without any cause. The use of graphs and implicit order constraints help to simplify the model while maintaining its properties.

Our representation for methods is very simplified and only gives causeless causal links. The causes of each link to the preconditions and effects of composite actions are inferred. This inference is done by running an instance of POP on the methods of each composite action *a* while compiling the domain. We use the following formula to compute the final preconditions and effects of *a*: $pre(a) = \bigcup_{a_s \in L^+(a)} causes(a_s)$ and $eff(a) = \bigcup_{a_s \in L^-(a)} causes(a_s)$. Errors are reported if POP cannot be completed or if nested methods contain their parent action as a step ($a \notin A_a^*$, see definition 10).

Problem

Problems instances are often most simply described by two components: the initial state and the goal.

Definition 5 (Problem). The planning problem is defined as a tuple $\mathcal{P} = \langle \mathcal{D}, C_p, \Omega \rangle$ where:

- \mathcal{D} is a planning domain.
- C_p is the set of **problem constants** disjoint from the domain constants.
- Ω is the **root operator** of the problem which methods are potential solutions of the problem.

Example: We use a simple problem for our example domain. The initial state provides that nothing is ready, taken or hot and all containers are empty (all using quantifiers). The goal is to have tea made. For reference, listing 2 contains the problem instance we use as an example.

```
1 init eff (hot(~), taken(~), placed(~), ~
    in ~);
2 goal pre (hot(water), tea in cup, water
    in cup, placed(spoon), placed(cup));
```

Listing 2: Example of a problem instance for the kitchen domain.

The root operator is initialized to $\Omega = \langle "", s_0, s^*, \{\pi_{lv(\Omega)}\} \rangle$, with s_0 being the initial state and s^* the goal specification. The method $\pi_{lv(\Omega)}$ is a partial order plan with the initial and goal steps linked together via Ω . The initial partial order plan is $\pi_{lv(\Omega)} = (\{I, G\}, \{I \xrightarrow{s_0} \Omega \xrightarrow{s^*} G\})$, with $I = \langle "init", \emptyset, s_0, \emptyset \rangle$ and $G = \langle "goal", s^*, \emptyset, \emptyset \rangle$.

Partial Order Planning

Our method is based on the classical POP algorithm. It works by refining a partial plan into a solution by recursively removing all of its flaws.

Definition 6 (Flaws). Flaws have a *proper fluent* *f* and a causing step often called the *needer* a_n . Flaws in a partial plan are either:

- **Subgoals**, *open conditions* that are yet to be supported by another step a_n often called *provider*. We note subgoals $\downarrow_f a_n$ (see definition 14).
- **Threats**, caused by steps that can break a causal link with their effects. They are called *breakers* of the threatened link. A step a_b threatens a causal link $l_t = a_p \xrightarrow{f} a_n$ if and only if $\neg f \in eff(a_b) \wedge a_p > a_b > a_n$. Said otherwise, the breaker can cancel an effect of a providing step a_p , before it gets used by its needer a_n . We note threats $a_b \boxtimes l_t$.

Example: Our initial plan contains two unsupported subgoals: one to make the tea ready and another to put sugar in it. In this case, the needer is the goal step and the proper fluents are each of its preconditions.

These flaws need to be fixed in order for the plan to be valid. In POP it is done by finding their resolvers.

Definition 7 (Resolvers). Classical resolvers are additional causal links that aim to fix a flaw.

- *For subgoals*, the resolvers are a set of potential causal links containing the proper fluent f in their causes while taking the needer step a_n as their target and a **provider** step a_p as their source.
- *For threats*, we usually consider only two resolvers: **demotion** ($a_b > a_p$) and **promotion** ($a_n > a_b$) of the breaker relative to the threatened link. We call the added causeless causal link a **guarding** link.

Example: The subgoal for $in(water, cup)$, in our example, can be solved by using the action $pour(water, cup)$ as the source of a causal link carrying the proper fluent as its only cause.

The application of a resolver does not necessarily mean progress. It can have consequences that may require reverting its application in order to be able to solve the problem.

Definition 8 (Side effects). Flaws that are caused by the application of a resolver are called *related flaws*. They are inserted into the *agenda*² with each application of a resolver:

- *Related subgoals* are all the new open conditions inserted by new steps.
- *Related threats* are the causal links threatened by the insertion of a new step or the deletion of a guarding link.

Flaws can also become irrelevant when a resolver is applied. It is always the case for the targeted flaw, but this can also affect other flaws. Those *invalidated flaws* are removed from the agenda upon detection:

- *Invalidated subgoals* are subgoals satisfied by the new causal links or the removal of their needer.
- *Invalidated threats* happen when the breaker no longer threatens the causal link because the order guards the threatened causal link or either of them have been removed.

Example: Adding the action $pour(water, cup)$ causes a related subgoal for each of the preconditions of the action which are: the cup and the water must be taken and water must not already be in the cup.

In algorithm 1 we present a generic version of POP inspired by Ghallab *et al.* (2004, sec. 5.4.2).

Example: The refinement process as illustrated in figure 1 follows the different steps of the POP algorithm. The first step is to find good candidates for resolution of the selected flaw. This is done in the `solve` function at line 5. Once a potentially compatible action is found, the algorithm search for a variable unification. If it is found the action is instantiated if needed. This gives as a result the list of all compatible resolvers. At line 7, the chosen resolver is applied and the action becomes a step of the plan. The last step of the refinement is to consider all side effects and invalidations caused by the resolver. This is done in line 8.

²An agenda is a flaw container used for the flaw selection of POP.

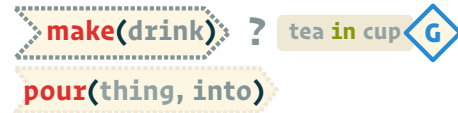
Algorithm 1 Partial Order Planner

```


1 function pop(Agenda  $a$ , Problem  $\mathcal{P}$ )
2   if  $a = \emptyset$  then  $\triangleright$  Populated agenda of flaws needs to be
   provided
3     return Success  $\triangleright$  Stops all recursion
4   Flaw  $f \leftarrow \text{choose}(a)$   $\triangleright$  Heuristically chosen flaw removed
   from agenda
5   Resolvers  $R \leftarrow \text{solve}(f, \mathcal{P})$ 
6   for all  $r \in R$  do  $\triangleright$  Non-deterministic choice operator
7     apply( $r, \pi$ )  $\triangleright$  Apply resolver to partial plan
8     Agenda  $a' \leftarrow \text{update}(a)$ 
9     if pop( $a', \mathcal{P}$ ) = Success then  $\triangleright$  Refining recursively
10      return Success
11     revert( $r, \pi$ )  $\triangleright$  Failure, undo resolver application
12    $a \leftarrow a \cup \{f\}$   $\triangleright$  Flaw wasn't resolved
13   return Failure  $\triangleright$  Revert to last non-deterministic choice

```

1. Resolver candidates



2. Variable unification



3. Resolver selection



4. Resolver application



5. Side effects search



Figure 1: The refinement process of our version of POP.

In definition 8, we mentioned effects that aren't present in classical POP, namely *negative resolvers*. All classical resolvers only add steps and causal links to the partial plan. Our method needs to remove composite steps and their adjacent links when expanding them.

The Heart Of The Method

In this section, we explain how our method combines POP and HTN and how they are used to generate intermediary abstract plans.

Additional Notions

In order to properly introduce the changes made for using HTN domains in POP, we need to define a few notions.

Definition 9 (Transposition). In order to transpose the causal links of an action a' with the ones of an existing step a , we do the following operation :

$$a \triangleright^- a' = \{\phi^-(l) \xrightarrow{\text{causes}(l)} a', l \in L_\pi^-(a)\} \cup (L_\pi^- \setminus L_\pi^-(a))$$

It is the same with $a' \xrightarrow{\text{causes}(l)} \phi^+(l)$ and L^+ for $a \triangleright^+ a'$. This supposes that the respective preconditions and effects of a and a' are equivalent. When not signed, the transposition is generalized: $a \triangleright a' = a \triangleright^- a' \cup a \triangleright^+ a'$.

Example: Using $a \triangleright^- a'$ means that all incoming links of a are now incoming links of a' instead.

Definition 10 (Proper actions). Proper actions are actions that are “contained” within an entity:

- For a *domain* or a *problem* it is $A_{\mathcal{D}|\mathcal{P}} = O$.
- For a *plan* it is $A_\pi = S_\pi$.
- For an *action* it is $A_a = \bigcup_{m \in \text{methods}(a)} S_m$.

Recursively: $A_a^n = \bigcup_{a' \in A_a} A_{a'}^{n-1}$.

We note $A_a^* = A_a^{lv(a)}$ the set of **extended proper actions** of the action a .

Example: The proper actions of $\text{make}(\text{drink})$ are the actions contained within its methods. The set of extended proper actions adds all proper actions of its single proper composite action $\text{infuse}(\text{drink}, \text{water}, \text{cup})$.

Definition 11 (Abstraction level). This is a measure of the maximum amount of abstraction an entity can express:³

$$lv(x) = \left(\max_{a \in A_x} (lv(a)) + 1 \right) [A_x \neq \emptyset]$$

Example: The abstraction level of any atomic action is 0 while it is 2 for the composite action $\text{make}(\text{drink})$. The example domain (in listing 1) has an abstraction level of 3.

³We use Iverson brackets here, see notations in table 1.

Abstraction In POP

The most straightforward way to handle abstraction with POP is illustrated in another planner called Duet (Gerevini *et al.* 2008) by managing hierarchical actions separately from a regular planner. We chose another way inspired by the works of Bechon *et al.* (2014) on a planner called HiPOP. The difference between the original HiPOP and our implementation of it is that we focus on the expressivity and the ways flaw selection can be exploited for partial resolution. Our version is lifted at runtime while the original is grounded for optimizations. All mechanisms we have implemented use POP but with different management of flaws and resolvers. The original algorithm 1 is left untouched.

One of those changes is that resolver selection needs to be altered for subgoals. Indeed, as stated by the authors of HiPOP : the planner must ensure the selection of high level operators in order to benefit from the hierarchical aspect of the domain, otherwise, adding operators only increases the branching factor.

We also need to add a way to deal with composite actions once inserted in the plan to reduce them to their atomic steps.

Definition 12 (Abstraction flaw). It occurs when a partial plan contains a non-atomic step. This step is the needer a_n of the flaw. We note it $a_n \oplus$.

- **Resolvers:** An abstraction flaw is solved with an **expansion resolver**. The resolver will replace the needer with one of its instantiated methods in the plan. This is done by linking all causal links to the initial and goal steps of the method as such: $a_n \triangleright^- I_m \wedge a_n \triangleright^+ G_m$, with $m \in \text{methods}(a_n)$.
- **Side effects:** An abstraction flaw can be related to the introduction of a composite action in the plan by any resolver and invalidated by its removal.

Example: When adding the step $\text{make}(\text{tea})$ in the plan to solve the subgoal that needs tea being made, we also introduce an abstraction flaw that will need this composite step replaced by its method using an expansion resolver. In order to the expansion of a composite action into a plan, all existing links are reported to the initial and goal step of the selected method, while the composite action and its links are removed from the plan.

The main differences between HiPOP and HEART in our implementations are the functions of flaw selection and the handling of the results (one plan for HiPOP and a plan per cycle for HEART). In HiPOP, the flaw selection is made by prioritizing the abstraction flaws. Bechon *et al.* (2014) state that it makes the full resolution faster. However, it also loses opportunities to obtain abstract plans in the process.

Cycles

The main focus of our work is toward obtaining **abstract plans** which are plans that are completed while still containing composite actions. In order to do that the flaw selection function will enforce cycles in the planning process.

Definition 13 (Cycle). A cycle is planning phase defined as a tuple $c = \langle lv(c), agenda, \pi_{lv(c)} \rangle$ where : $lv(c)$ is the maximum abstraction level allowed for flaw selection in the *agenda* of

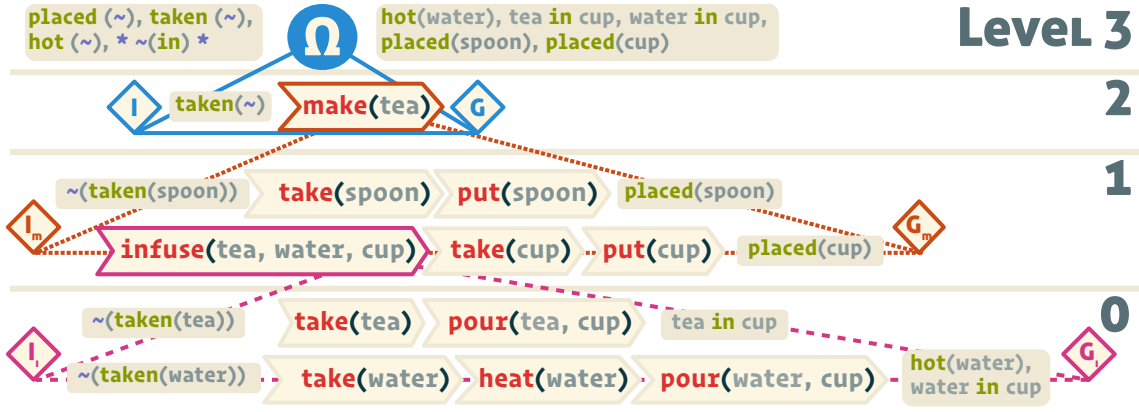


Figure 2: The cycle process on the example domain. Atomic actions that are copied from a cycle to the next are omitted.

remaining flaws in partial plan $\pi_{lv(c)}$. Resolvers of subgoals are therefore constrained by the following: $a_p \downarrow_f a_n : lv(a_p) \leq lv(c)$.

During a cycle all abstraction flaws are delayed. Once no more flaws other than abstraction flaws are present in the agenda, the current plan is saved and all remaining abstraction flaws are solved at once before the abstraction level is lowered for the next cycle: $lv(c') = lv(c) - 1$. Each cycle produces a more detailed abstract plan than the one before.

Abstract plans allow the planner to do an approximative form of anytime execution. At any given time the planner is able to return a fully supported plan. Before the first cycle, the plan returned is $\pi_{lv(\Omega)}$.

Example: In our case using the method of intent recognition of Sohrabi *et al.* (2016), we can already use $\pi_{lv(\Omega)}$ to find a likely goal explaining an observation (a set of temporally ordered fluents). That can make an early assessment of the probability of each goal of the recognition problem.

For each cycle c , a new plan $\pi_{lv(c)}$ is created as a new method of the root operator Ω . These intermediary plans are not solutions of the problem, nor do they mean that the problem is solvable. In order to find a solution, the HEART planner needs to reach the final cycle c_0 with an abstraction level $lv(c_0) = 0$. However, these plans can be used to derive meaning from the potential solution of the current problem and give a good approximation of the final result before its completion.

Example: In the figure 2, we illustrate the way our example problem is progressively solved. Before the first cycle c_2 all we have is the root operator and its plan π_3 . Then within the first cycle, we select the composite action *make(tea)* instantiated from the operator *make(drink)* along with its methods. All related flaws are fixed until all that is left in the agenda is the abstract flaws. We save the partial plan π_2 for this cycle and expand *make(tea)* into a copy of the current plan π_1 for the next cycle. The solution of the problem will be stored in π_0 once found.

Properties

Soundness

For an algorithm to be sound, it needs to provide only *valid* solutions. In order to prove soundness, we first need to define the notion of support.

Definition 14 (Support). An open condition f of a step a is supported in a partial order plan π if and only if

$$\exists l \in L_{\pi}^-(a) \wedge \nexists a_b \in S_{\pi} : f \in \text{causes}(l) \wedge (\phi^-(l) > a_b > a \wedge \neg f \in \text{eff}(a_b))$$

This means that the fluent is provided by a causal link and isn't threatened by another step. We note support $\pi \downarrow_f a$.

Full support of a step is achieved when all its preconditions are supported: $\pi \Downarrow a \equiv \forall f \in \text{pre}(a) : \pi \downarrow_f a$.

We also need to define validity in order to derive all its logical equivalences for the proofs:

Definition 15 (Validity). A plan π is a valid solution of a problem \mathcal{P} if and only if $\forall a \in S_{\pi} : \pi \Downarrow a \wedge lv(a) = 0$.

We can now start to prove the soundness of our method. We base this proof upon the one done in (Erol *et al.* 1994). It states that for classical POP if a plan doesn't contain any flaws, it is fully supported. Our main difference being with abstraction flaws we need to prove that its resolution doesn't leave classical flaws unsolved in the resulting plan.

Lemma (Expansion with an empty method). *When a fully supported composite action $\pi \Downarrow a$ is expanded using an empty method $m = (\{I_m, G_m\}, \{I_m \rightarrow G_m\})$, adding all open conditions of G_m as subgoals will result in a plan π' , without any undiscovered flaws.*

Proof. The initial and goal steps of a method follows ($\text{pre}(a) = \text{eff}(a)$). By definition of full support:

$$L_{\pi'}^-(I_m) = L_{\pi}^-(a) \wedge \text{pre}(I_m) = \text{pre}(a) \implies (\pi \Downarrow a \equiv \pi' \Downarrow I_m)$$

The only remaining undiscovered open conditions in the plan π' are therefore those caused by G_m : $\{f \in \text{pre}(G_m) : \pi' \nmid_f G_m\}$.

No new threats are introduced since the link between I_m and G_m is causeless and they inherit the order of the parent composite action. All added actions are atomic so no abstraction flaws are added.

□

Lemma (Expansion with an arbitrary method). *If a fully supported composite action $\pi \Downarrow a$ is replaced by an arbitrary method m and all open conditions contained within S_m as subgoals and all threatened links within L_m as threats are added to the agenda, the resulting plan π' will not have any undiscovered flaws.*

Proof. When replacing a composite action with a method in an existing plan we do the following operations: $S_{\pi'} = S_m \cup (S_\pi \setminus a)$, $L_{\pi'}(I_m) = a \triangleright I_m$. The only added flaws are then:

$$\bigcup_{a_m \in S_m} \pi' \downarrow_f a_m \quad \bigcup_{a_b \in S_{\pi'}} l \in L_{\pi'} a_b \otimes l \quad \bigcup_{a_c \in S_m} lv(a_c) \neq 0$$

That means that all added actions have their subgoals considered and all links have their threats taken into account along with all additional abstraction flaws.

This proves that expansion does not introduce flaws that are not added to the POP agenda. Since POP must resolve all flaws in order to be successful and according to the proof of the soundness of POP, HEART is sound as well.

□

Another proven property is that intermediary plans are valid in the classical definition of the term (without considering abstraction flaws) and when using this definition, HEART is sound on its anytime results too.

Completeness

The completeness of POP has been proven in the same paper as for its soundness (Erol *et al.* 1994). Since our method uses the same algorithm only the differences must be proven to respect the contract of application and reversion. This contract states that applying a resolver prevents the recurrence of its flaw and that reverting the application of a resolver must restore the plan and agenda to their previous state.

Lemma (Expansion solves the abstraction flaw). *The application of an expansion resolver invalidates the related abstraction flaw.*

Proof. Abstraction flaws arise from the existence of their related composite step a in the plan. Since the application of an expansion resolver is of the form $S'_\pi = (S_\pi \setminus a) \cup S_m$ unless $a \in S_m$ (which is forbidden by definition 3), therefore $a \notin S'_\pi$.

□

Lemma (Solved abstraction flaws cannot reoccur). *The application of an expansion resolver on a plan π , guarantees that $a \notin S_\pi$ for any partial plan refined from π without reverting the application of the resolver.*

Proof. As stated in the definition of the methods (definition 3): $a \notin A_a^*$. This means that a cannot be introduced in the plan by its expansion or the expansion of its proper actions. Indeed, once a is expanded, the level of the following cycle $c_{lv(a)-1}$ prevents a to be selected by subgoal resolvers. It cannot either be contained in the methods of another action that are selected afterward because otherwise following definition 11 its level would be at least $lv(a) + 1$.

Since the implementation guarantees that the reversion is always done without side effects, all the aspects of completeness of POP are preserved in HEART.

□

Results

In order to assess its capabilities, HEART was tested on two criteria: quality and complexity. All tests were executed on an Intel® Core™ i7-7700HQ CPU clocked at 2.80GHz. The Java process used only one core and wasn't limited by time or memory. Each experiment was repeated between 700 and 10 000 times to ensure that variations in speed weren't impacting the results.




Figure 3: Evolution of the quality with computation time.

figure 3 shows how the quality is affected by the abstraction in partial plans. The tests are made using our example domain (see listing 1). The quality is measured by counting the number of providing fluents in the plan $|\bigcup_{a \in S_\pi} eff(a)|$, which is actually used to compute the probability of a goal in intent recognition. These results show that in some cases it may be more interesting to plan in a leveled fashion to solve HTN problems. For the first cycle of level 3, the quality of the abstract plan is already of 60%. This is the quality of the exploitation of the plan *before any planning*. With almost 3 quarters of the final quality and less than half of the complete computation time, the result of the first cycle is a good quality/time compromise.

In the second test, we used generated domains. These domains consist of an action of abstraction level 5. It has a single method containing a number of actions of level 4. We call this number the width of the domain. All needed actions are built recursively to form a tree shape. Atomic actions only have single fluent effects. The goal is the effect of the higher level action and the initial state is empty. These domains do not contain negative effects. figure 4 shows the computational profile of HEART for various levels and widths. We note that the behavior of HEART seems to follow an exponential law with the negative exponent of




Figure 4: Impact of domain shape on the computation time by levels. The scale of the vertical axis is logarithmic. Equations are the definition of the trend curves.

the trend curves seemingly being correlated to the actual width. This means that computing the first cycles has a complexity that is close to being *linear* while computing the last cycles is of the same complexity as classical planning which is at least *P-SPACE* (Weld 1994).

Conclusions

In this paper, we have presented a new planner called HEART based on POP. We showed how HEART performs compared to complete planners in terms of speed and quality. While the abstract plans generated during the planning process are not complete solutions, they are exponentially faster to generate while retaining significant quality over the final plans. By using these plans it is possible to find good approximations to intractable problems within tight time constraints.

References

- Baiocchi, M., S. Marcugini, and A. Milani
Encoding planning constraints into partial order planning domains, *International Conference on Principles of Knowledge Representation and Reasoning*, Morgan Kaufmann Publishers Inc., 1998, 608–616.
- Baker, C. L., J. B. Tenenbaum, and R. R. Saxe
Goal inference as inverse planning, *Proceedings of the Annual Meeting of the Cognitive Science Society*, Vol. 29, 2007.
- Bechon, P., M. Barbier, G. Infantes, C. Lesire, and V. Vidal
HiPOP: Hierarchical Partial-Order Planning, *European Starting AI Researcher Symposium*, IOS Press, 2014, 264,51–60.
- Chen, J., Y. Chen, Y. Xu, R. Huang, and Z. Chen
A Planning Approach to the Recognition of Multiple Goals, *IJIS*, 28 (3), 203–216, 2013. doi:[10.1002/int.21565](https://doi.org/10.1002/int.21565).
- Erol, K., J. A. Hendler, and D. S. Nau
UMCP: A Sound and Complete Procedure for Hierarchical Task-network Planning, *Proceedings of the International Conference on Artificial Intelligence Planning Systems*, University of Chicago, Chicago, Illinois, USA: AAAI Press, June 1994, 2,249–254.
- Fox, M., D. Long, and D. Magazzeni
Explainable Planning, *Proceedings of IJCAI Workshop on Explainable AI*, Melbourne, Australia, August 2017.
- Gerevini, A., U. Kuter, D. S. Nau, A. Saetti, and N. Waisbrot
Combining Domain-Independent Planning and HTN Planning: The Duet Planner, *Proceedings of the European Conference on Artificial Intelligence*, 2008, 18,573–577.
- Ghallab, M., D. Nau, and P. Traverso
Automated planning: Theory & practice, Elsevier, 2004.
- Göbelbecker, M., T. Keller, P. Eyerich, M. Brenner, and B. Nebel
Coming Up With Good Excuses: What to do When no Plan Can be Found, *Proceedings of the International Conference on International Conference on Automated Planning and Scheduling*, AAAI Press, May 2010, 20,81–88.
- Höller, D., P. Bercher, G. Behnke, and S. Biundo
Plan and Goal Recognition as HTN Planning, *Proceedings of AAAI Workshop on Plan, Activity and Intent Recognition*, 2018.
- Pellier, D., and H. Fiorino
A unified framework based on HTN and POP approaches for multi-agent planning, *Proceedings of the International Conference on Intelligent Agent Technology*, IEEE/wic/acm; IEEE Computer Society, 2007, 285–288.
- Ramirez, M., and H. Geffner
Plan recognition as planning, *Proceedings of the International Conference on International Conference on Automated Planning and Scheduling*, AAAI Press, 2009, 19,1778–1783.
- Ramirez, M., and H. Geffner
Probabilistic plan recognition using off-the-shelf classical planners, *Proceedings of the Conference of the Association for the Advancement of Artificial Intelligence*, 2010, 24,1121–1126.
- Sohrabi, S., A. V. Riabov, and O. Udrea
Plan Recognition as Planning Revisited, *Proceedings of the International Joint Conference on Artificial Intelligence*, Vol. 25, 2016.
- Talamadupula, K., G. Briggs, T. Chakraborti, M. Scheutz, and S. Kambhampati
Coordination in human-robot teams using mental modeling and plan recognition, *Intelligent Robots and Systems*, Chicago, IL, USA: IEEE, September 2014, 2957–2962. doi:[10.1109/IROS.2014.6942970](https://doi.org/10.1109/IROS.2014.6942970).
- Weld, D. S.
An introduction to least commitment planning, *AI Mag.*, 15 (4), 27, 1994.